

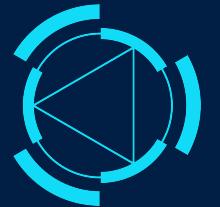


RDFox Reasoning Workshop



The world's most performant knowledge graph and semantic reasoning engine.

Requirements



A. Get an RDFox License

<https://www.oxfordsemantic.tech/tryrdfoxforfree>

B. Download RDFox (& unzip)

<https://www.oxfordsemantic.tech/downloads>

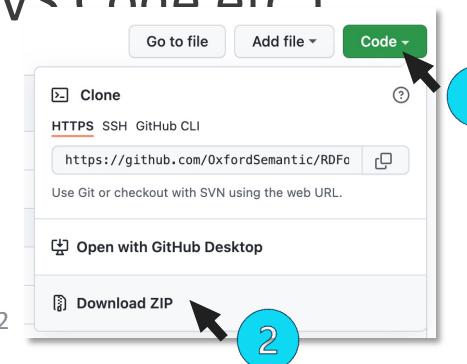
C. Download the class materials from Github:

<git@github.com:OxfordSemantic/RDFoxWorkshop-Reasoning.git>

D. OPTIONAL Get your IDE of choice ready (VS Code etc.)

<https://code.visualstudio.com/>

The screenshot shows two pages. The top part is a 'Free 30 day trial' form asking for a short questionnaire to request a license. The bottom part is a 'Downloads' page for RDFox v5.5, showing download links for various operating systems: MAC (INTEL), MAC (ARM), LINUX (x86), LINUX (ARM), and WINDOWS.



You should have...



examples	✓	30/03/2022 11:23	File folder
include	✓	30/03/2022 11:23	File folder
lib	✓	30/03/2022 11:23	File folder
RDFox.lic	✓	30/03/2022 11:23	Text Document
RDFox	✓	30/03/2022 11:23	Application
readme	✓	30/03/2022 11:23	Text Document
version	✓	30/03/2022 11:23	Text Document



Make sure to put your license
INSIDE the **RDFox folder**.

axioms	✓	26/04/2022 16:48	File folder
data	✓	26/04/2022 16:48	File folder
queries	✓	26/04/2022 16:48	File folder
rules	✓	26/04/2022 16:48	File folder
.gitignore	✓	26/04/2022 16:48	Git Ignore Source ... 1 KB
LICENSE	✓	26/04/2022 16:48	File 2 KB
protege-ontology	✓	26/04/2022 16:48	TTL File
RDFoxWorkshop-Reasoning	✓	26/04/2022 16:48	Chrono
start-workshop.rdfox	✓	26/04/2022 16:48	RDFox



Open the class PDF so you can follow along!

You will need to follow the links provided.

Objectives



By the end of the class everyone will:

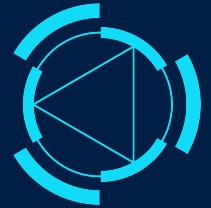
- know how to write OWL axioms in Turtle syntax
- know how to add OWL axioms to RDFox datastores
- know how to write and import Datalog rules
- understand how RDFox performs reasoning
- know common uses of rules

Workshop Structure



- The repository contains an RDFox script file “start-workshop.rdf0x” with all the commands you need – copy and paste them one by one as we go
- Each axiom/rule has an associated query (mentioned on the slide). Run the query before and after reasoning to see the difference!
- Some examples are exercises where you will need to fill in the gaps
- Answers to the exercises have “completed” in their names

Workshop Structure

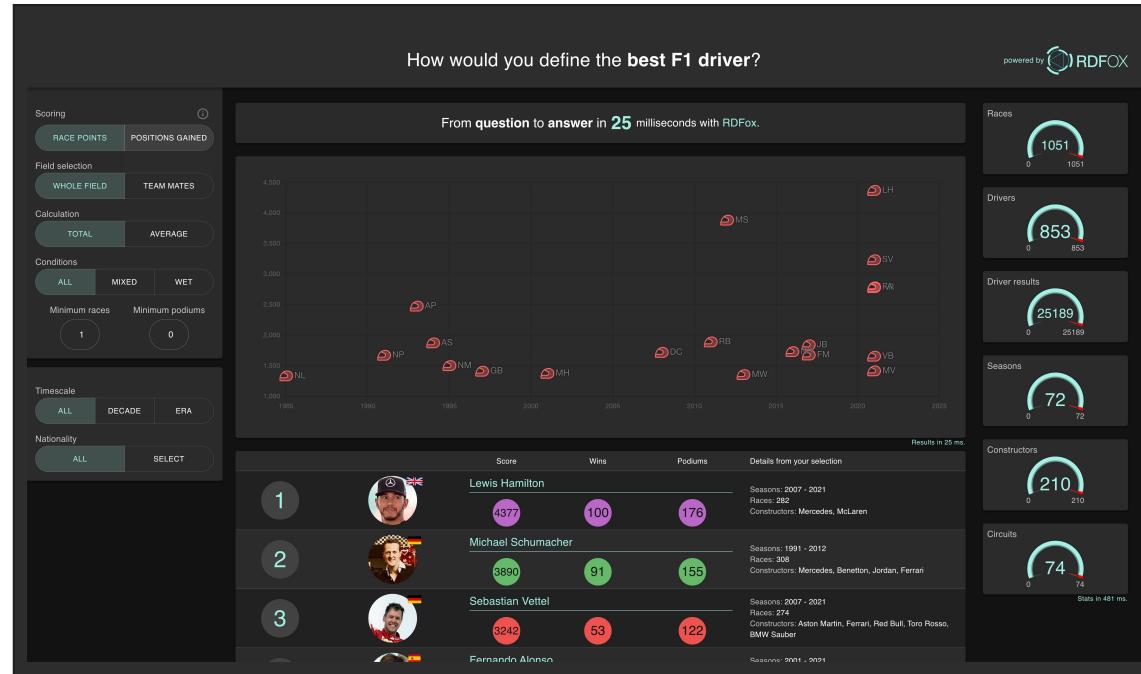


- I will talk about a concept, what it is, and why it's relevant
- I will go through a set of commands to show you this
- You are encouraged to follow along
- If there is an exercise, I will give you 1-2 minutes to try this, then I will go through it

Who is the Greatest Formula One Driver of All Time?



Controls and filters for the scoring system.



Statistics about the data used to form the results.

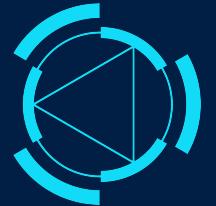
This matches the filters.

Try it for yourself!

<http://f1.rdfbox.tech>

© Oxford Semantic Technologies 2022

Setting up RDFox



- We recommend using an IDE (e.g. VS Code)
- Open a terminal, navigate to the workshop folder
`cd <path_to_workshop_folder>`
and from there run
`<path_to_RDFox>/RDFox sandbox`
(Mac) or
`<path_to_RDFox>\RDFox.exe sandbox`
(Windows)
- RDFox server should be running now:

```
Source code for RDFox v1.0 Copyright 2013 Oxford University Innovation Limited and subsequent improvements Copyright 2017-2021 by Oxford Semantic Technologies Limited.
```

```
This copy of RDFox is licensed for Developer use to Tom Vout (tom.vout@oxfordsemantic.tech) of OST until 07-Jun-2022 16:01:44
```

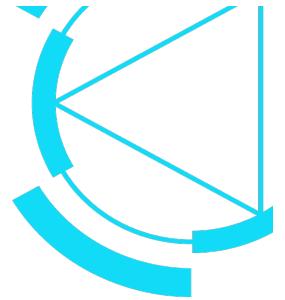
```
This system is equipped with 16.9 GB of RAM, and RDFox is configured to use at most 15.2 GB (89.9% of the total).
```

```
Currently, 2.8 GB (18.4% of the amount allocated to RDFox) appear to be available on the system.
```

```
Since RDFox is a RAM-based system, its performance can suffer when other running processes use a lot of memory.
```

```
A new server connection was opened as role 'guest' and stored with name 'sc1'.
```

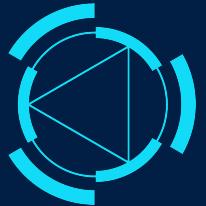
```
> []
```



OWL reasoning

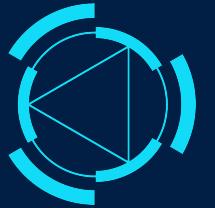
- Intro
- Subclasses
- Property types and domain/range
- Fixing mistakes
- Restrictions
- Open vs Closed World
- Property Chain
- A more complex example

Web Ontology Language (OWL)

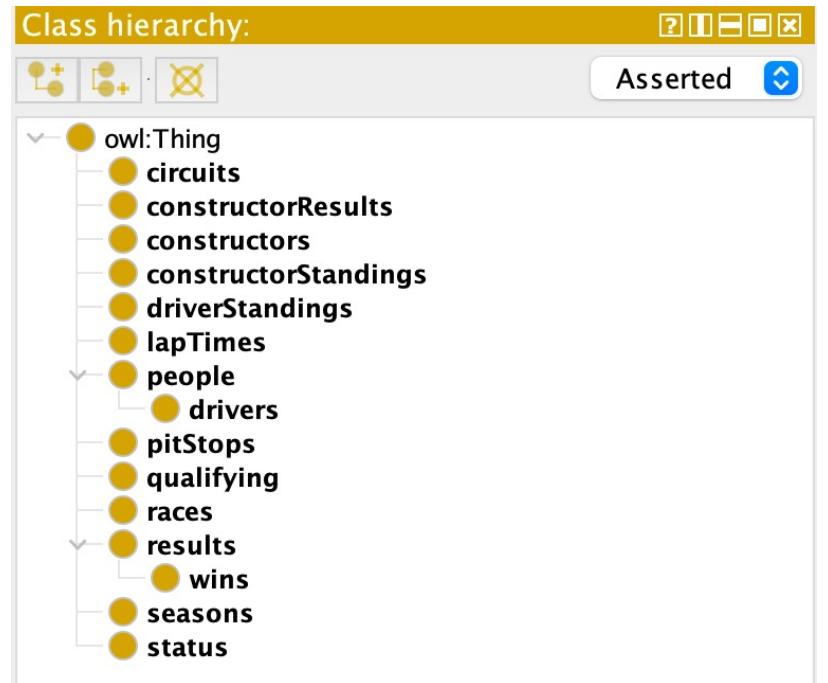


- standard created by W3C, RDFox supports the profile OWL 2 RL
- axioms = assertions about classes, properties or individuals
- ontology = collection of axioms
- ontologies are themselves graphs and can be written in many different syntaxes - we will use Turtle
- internally, RDFox translates axioms into Datalog rules

Protégé



- useful program for managing ontologies
- free and open source
- before writing an ontology, it is often better to have some data first



Sample data



- we import real F1 race data from file “upTo2020.ttl”
- we also import sample data suited to showing how OWL axioms work
(file “sample.ttl”)

```
:drivers\123456789 # no class
|   :drivers_forename "John";
|   :drivers_surname "Smith" .

:results\123456789 a :results;
|   :results_driver :drivers\123456789, :drivers\1; # two drivers
|   :drivers_forename "John"; # wrong – this is a driver property
|   # no constructor
|   :results_race :races\123456789;
|   :results_positionOrder "1"^^xsd:integer .
```

Loading Data into RDFox



First we need to create a data store...

```
dstore create f1
```

Then to set it as active.

```
active f1
```

Now import the race data up to the year 2020.

```
import data/upTo2020.ttl
```

And finally import our error-filled sample data.

```
import data/sample.ttl
```

Axiom set 1

Subclasses



```
@prefix : <http://www.oxfordsemantic.tech/f1demo/> .  
  
:drivers a owl:Class .  
:people a owl:Class .  
  
:drivers rdfs:subClassOf :people .
```

this is a completely new class that was not previously present in the data

- declares that `:drivers` and `:people` are classes
- says that `:drivers` is a subclass of `:people`
- effect: all instances of `:drivers` also become instances of `:people`
(check by running query 1 – q1.rq)

Axiom set 2

Property types and domain/range



```
@prefix : <http://www.oxfordsemantic.tech/f1demo/> .  
  
:drivers a owl:Class .  
  
:drivers_forename a owl:DatatypeProperty;  
|   rdfs:domain :drivers .
```

OWL defines each property as either a “data” or “object” property

- says the domain of the property `:drivers_forename` is the class `:drivers`
- effect: `:drivers\123456789` and `:results\123456789` from sample both become instances of `:drivers` (check by running query 2 – q2.rq)

Fixing Mistakes



What should I do if I make a mistake?

```
import axioms :axioms -
```

```
update ! clear graph :axioms
```

```
import > :axioms axioms/a1.ttl axioms/a2.ttl ...
```

```
import axioms :axioms
```

Axiom set 3 - exercise

Property types and domain/range



```
@prefix : <http://www.oxfordsemantic.tech/f1demo/> .  
  
____ a owl:Class .  
____ a owl:Class .  
  
:results_race a owl:ObjectProperty ;  
|   rdfs:domain ____ ;  
|   rdfs:range ____ .
```

- establishes the domain and range of the `:results_race` property
- effect: `:races\123456789` from the sample becomes an instance of `:races` (check by running query 3 – q3.rq)

Axiom set 4

Restrictions

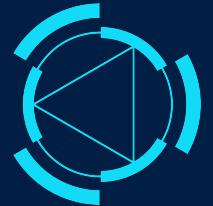


```
@prefix : <http://www.oxfordsemantic.tech/f1demo/> .  
  
:results_driver a owl:ObjectProperty .  
  
:results rdfs:subClassOf [  
    a owl:Restriction ;  
    owl:onProperty :results_driver ;  
    owl:maxCardinality 1  
] .
```

- says that every node of class `:results` has at most one value of `:results_driver`
- effect: property `owl:sameAs` now connects `:driver\123456789` and `:driver\1`, as well as every driver who ever raced to themselves (check by running query 4 – q4.rq)

Axiom set 5 - exercise

Restrictions



```
@prefix : <http://www.oxfordsemantic.tech/f1demo/> .  
  
:results_constructor a __ .  
  
  __ rdfs:subClassOf [  
    a __ ;  
    __ :results_constructor ;  
    owl:minCardinality __  
  ] .
```

- says that every node of class `:results` has at least one value of `:results_constructor`
- effect: none (Why? Because of the open world assumption of OWL)

Axiom set 6

Property chains



```
@prefix : <http://www.oxfordsemantic.tech/f1demo/> .  
  
:results_driver a owl:ObjectProperty .  
  
:hasRacedIn owl:propertyChainAxiom (  
    [owl:inverseOf :results_driver]  
    :results_race  
) .
```

- defines a new property `:hasRacedIn` as a combination of `:results_driver` (inverted) and `:results_race`
- effect: property `:hasRacedIn` now connects drivers and their races (check by running query 6 – q6.rq)

Axiom set 7

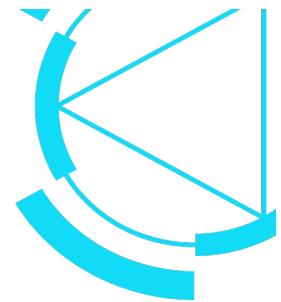
A more complex example



```
@prefix : <http://www.oxfordsemantic.tech/f1demo/> .  
  
:qualifying a owl:Class .  
  
:qualifying_position a owl:DatatypeProperty .  
  
:polePositionQualifying a owl:Class ;  
    owl:equivalentClass [  
        a owl:Class ;  
        owl:intersectionOf (  
            :qualifying  
            [  
                a owl:Restriction ;  
                owl:onProperty :qualifying_position ;  
                owl:hasValue "1"^^xsd:integer  
            ]  
        )  
    ].
```

- says that the **:polePositionQualifying** class is the intersection of the **:qualifying** class and nodes that have 1 as the value of their **:qualifying_position** property
- effect: some **:qualifying** instances also become instances of **:polePositionQualifying** (check by running query 7 – q7.rq)

Datalog Reasoning with RDFox



- Basic rules
- Filters
- Aggregates
- Negation
- Binds
- Incremental reasoning

Datalog Rules



$[fact A] :- [fact B]$

“*A is true whenever B is true*”

Rule 1

Basic rules



```
[?driver, :hasRacedIn, ?race] :-  
    [?result, :results_race, ?race],  
    [?result, :results_driver, ?driver] .
```

This rule will add a direct link between drivers and the races they raced in. Notice this has the same effect as axiom 6, but expressed more clearly

The same could seemingly be achieved with an equivalent SPARQL update, but rules offer several advantages over write queries.

Check for inferred triples with query 8 – q8.rq

Rule 2

Filters



```
[?driver, :hasPodiumInRace, ?race] :-  
    [?result, :results_race, ?race],  
    [?result, :results_driver, ?driver],  
    [?result, :results_positionOrder, ?positionOrder],  
    FILTER(?positionOrder < 4) .
```

We can use **FILTERs** in rules too, in this case to find the races in which a driver reached the podium.

Check for inferred triples with query 9 – q9.rq

Rule 3

Aggregates



```
[?driver, :hasRaceCount, ?raceCount] :-  
    AGGREGATE(  
        [?driver, :hasRacedIn, ?race]  
    ON ?driver  
    BIND COUNT(?race) AS ?raceCount  
    ) .
```

This rule adds a count of races a driver has entered.

Notice we use the previously inferred `:hasRacedIn` property. Rules can be composed together, and RDFox will automatically find the order in which to run them.

Check for inferred triples with query 10– q10.rq

Rule 4 - exercise

Aggregates



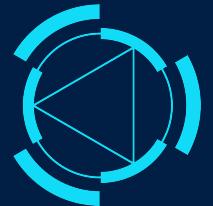
```
[?__, :hasRaceWinCount, ?__] :-  
    AGGREGATE(  
        [?__, :results_driver, ?__],  
        [?__, :results_positionOrder, 1]  
    ON ?driver  
    BIND COUNT(?__) AS ?__  
    ) .
```

This rule adds a count of races a driver has *won* (provided they won at least 1). We can put more than one atom *inside* the **AGGREGATE** statement.

Check for inferred triples with query 11 – q11.rq

Rule 5

Negation

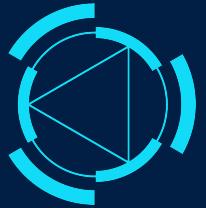


```
[?driver, a, :DriverWithoutPodiums] :-  
    [?driver, a, :drivers],  
    NOT EXISTS ?race IN (  
        | [?driver, :hasPodiumInRace, ?race]) .
```

This rule tells us that if a driver does not have a race where they were on the podium, then they are a :DriverWithoutPodiums.

So, we are looking for something that is **not** in the data (i.e. no race where the driver was on the podium).

Check for inferred triples with query 12 – q12.rq



Negation as Failure

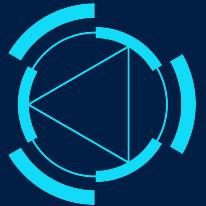
Negation as failure is negation understood as the *absence* of a triple.

Using Negation as Failure leads to ‘non-monotonic reasoning’- the engine may need to retract some triples instead of just adding them. This adds another dimension to an already complicated problem of reasoning planning, but RDFox handles all of that automatically.

Example:

What if a driver without any podiums *so far* gets a podium next year? In 2021, for instance, George Russell got his first podium. In our data, which is accurate up to and including 2020, George Russell would be a :DriverWithoutPodiums. But if we add data from 2021, we will need to retract this fact.

- Exercise: Run the query “q12b.rq”, then import the data from file “2021-22.ttl” and try again



Rule 6 - exercise

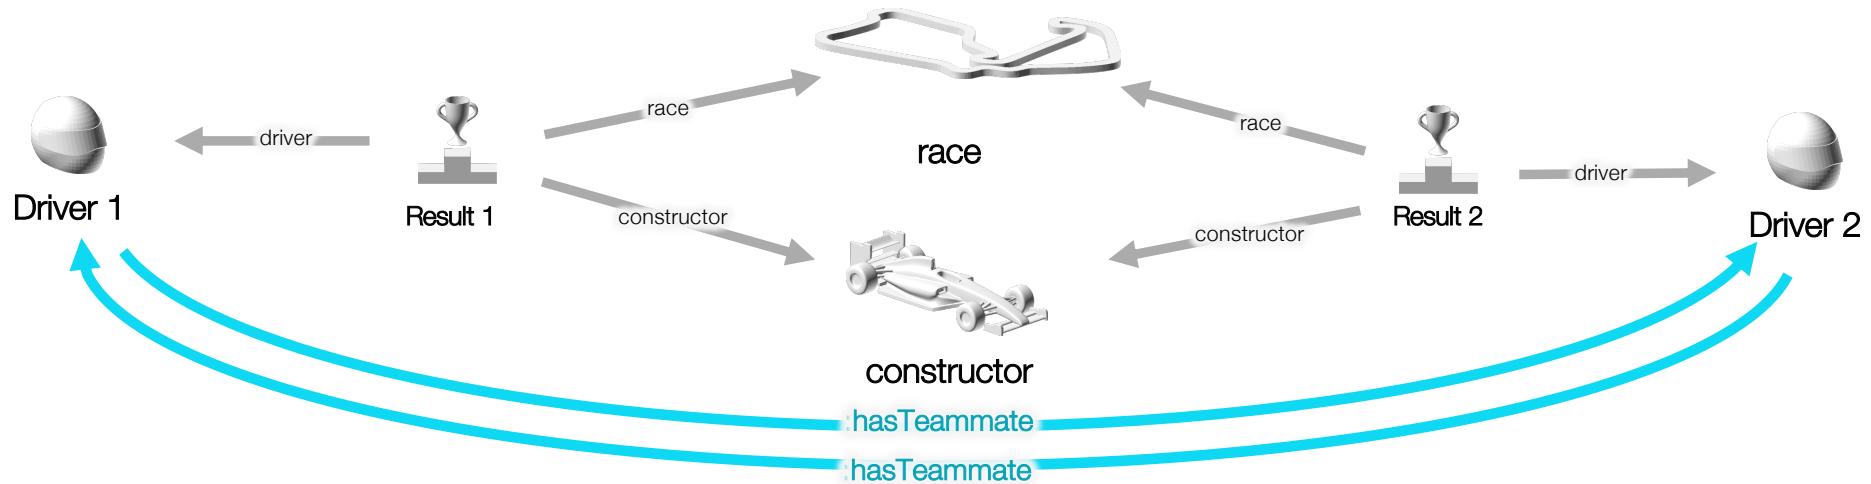
Binds

```
[?driver, :hasWinPercentage, ?percentage] :-  
    | ?___, :hasRaceCount, ?__],  
    | ?___, :hasRaceWinCount, ?__],  
    | BIND(?raceWinCount/?raceCount AS ?percentage) .
```

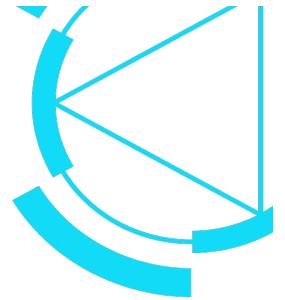
With **BIND** we can use various mathematical functions, string manipulation, regular expression matching, conditional binds, hashing and IRI creation.

Check for inferred triples with query 13 – q13.rq

Bonus Exercise

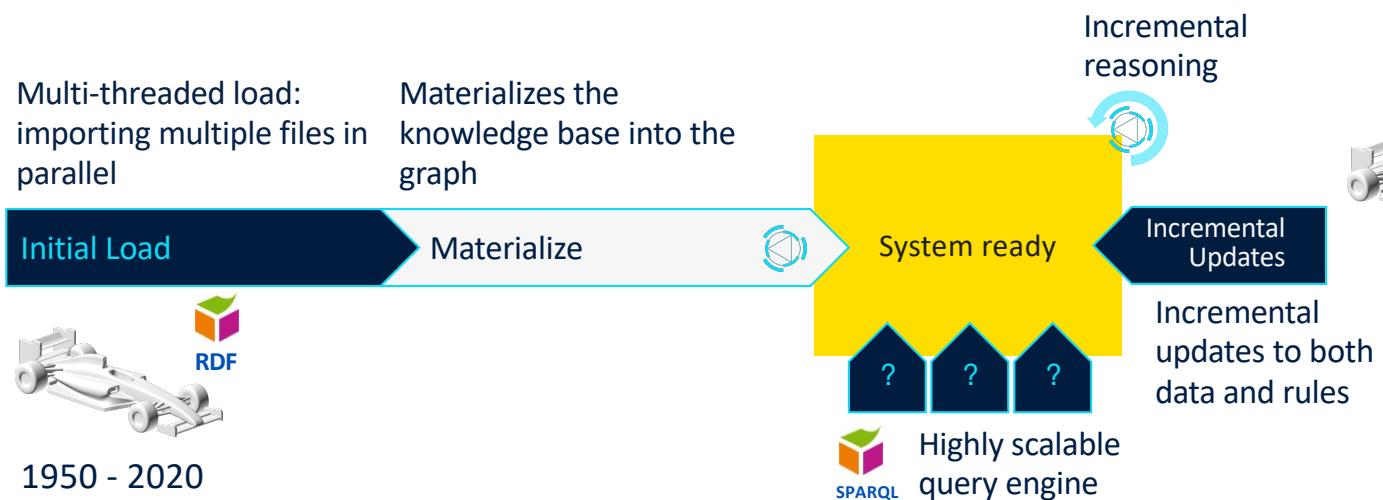


- Write and import a rule to find out if two drivers have been teammates (use the predicate **:hasTeammate**).



RDFox: How does it work?

Incremental Reasoning



Incremental Reasoning



- One of the key strengths of RDFox
- Data can be imported at any time, and the RDFox will automatically compute the necessary inferences
- This all happens *incrementally*, i.e. without the need to reboot.
- Incremental reasoning open the door to many novel use cases which were previously impractical

RDFox vs. Other Solutions



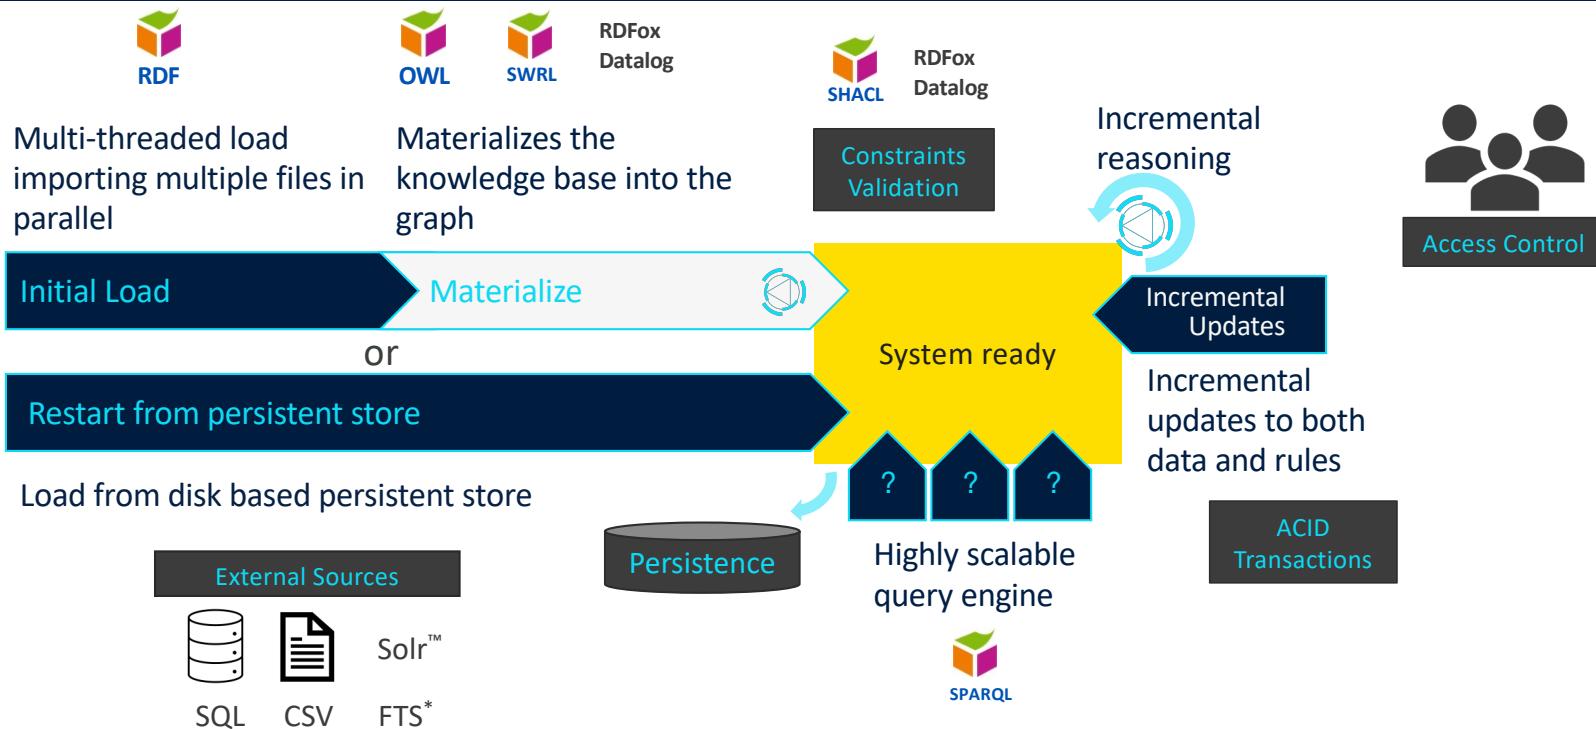
Feature	Example	RDFox	Materialisation Competitors	Query Rewriting Competitors
	"List all financial instruments / derivative instruments / futures." [class inferencing]	✓	✓*	✓**
	"What is the nearest electrical switch for a circuit?" [negation + recursion]	✓		
	"What is the total volume and value of trades?" [aggregation + arithmetic]	✓		
	"No trades over a limit & no traders without trades!" [filters + negation]	✓	Partial (SHACL)	
	"Every component must be connected to a power source!" [negation + recursion]	✓		
	"What is the total value of assets owned through holdings?"	✓		

* many times slower than RDFox

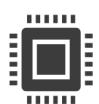
** orders of magnitude slower than RDFox

© Oxford Semantic Technologies

How does it work?



Deployment Options



Edge and mobile devices



Personal machines, on-premises to cloud

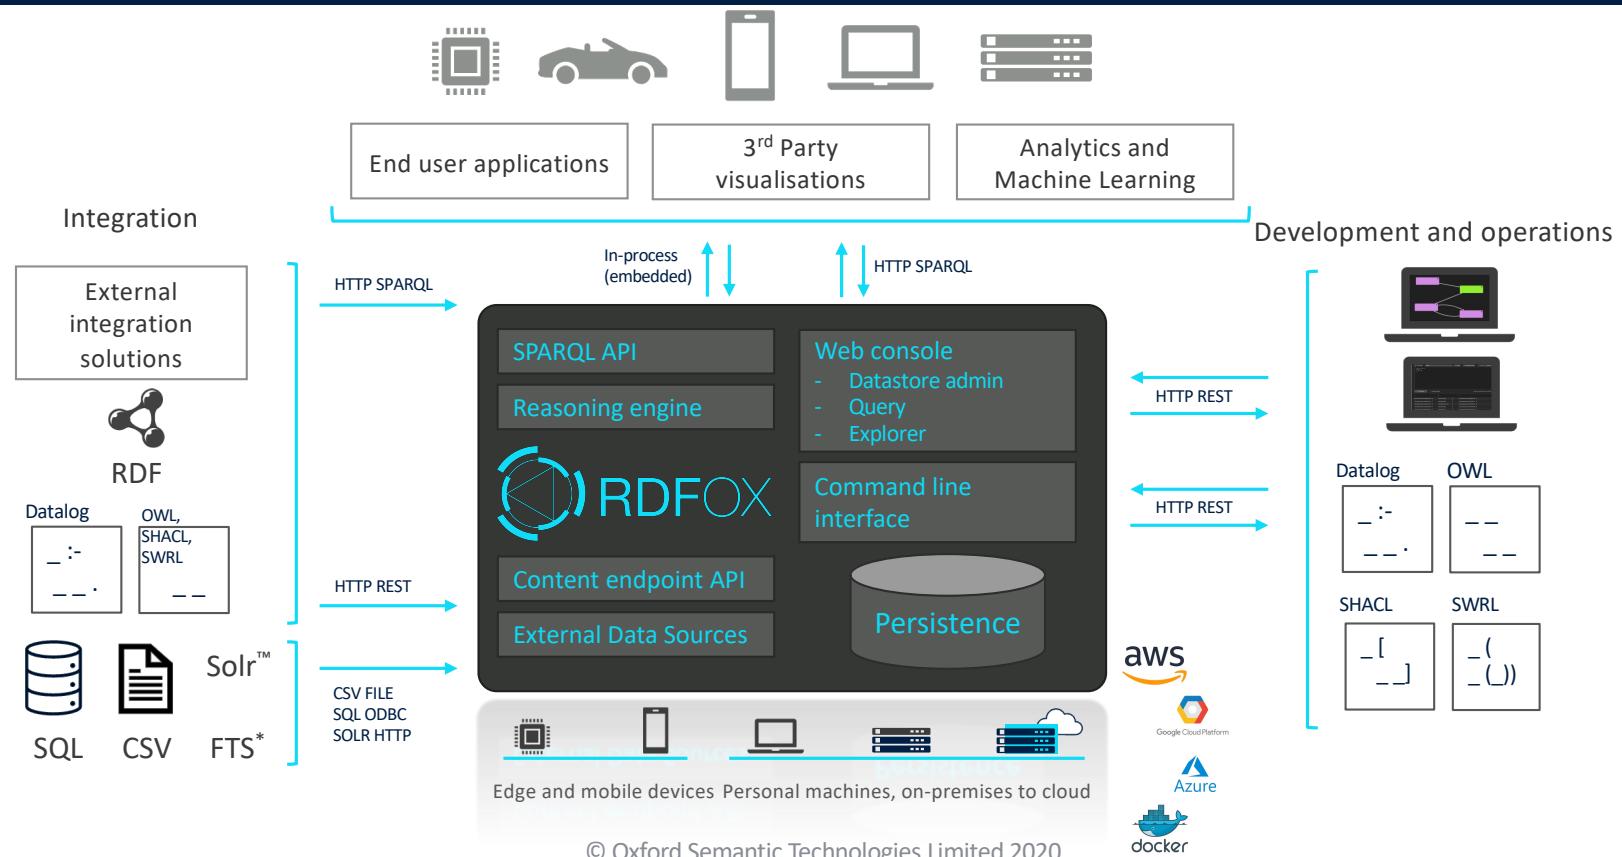


Google Cloud Platform



Architecture

RDFox system context diagram



Your final exercise...



Click on the link below to start filling out our quick survey. It won't take you more than 3 minutes!

<https://oxfordsemantictech.typeform.com/to/fygpYBrS>